

Creating a Linux virtual machine cluster and running simulation analysis with IPython Cluster

Overview

In this hands-on lab you will learn to deploy an IPython notebook cluster on Microsoft Azure. You will execute the Monte Carlo simulation example with IPython Notebook on the cluster you created.

Objectives

In this hands-on lab, you will learn how to:

- Provision a Linux virtual machine cluster with Python.
- Deploy IPython notebook on your virtual machines.
- Run Monte Carlo Simulation on IPython Notebook on Microsoft Azure.

Prerequisites

The following is required to complete this hands-on lab:

- A Microsoft Azure subscription - [sign up for a free trial](#)
 - Lab: Using Microsoft Azure Virtual Machine
-

Exercises

This hands-on lab includes the following exercises:

1. [Build an Environment to manage Microsoft Azure with Python](#)
2. [Deploy IPython notebook on Microsoft Azure](#)
3. [Run Monte Carlo Simulation on IPython](#)

Estimated time to complete this lab: **60** minutes.

Exercise 1: Build an Environment to manage Microsoft Azure with Python

First, you will need to deploy required software on your Linux machine. You are going to install git, Python 2.7, workerpool and paramiko, then you will connect Microsoft Azure by Python with some configuration. For the lab, we will either provide you with a shared virtual machine to do the exercise or let you provision a Linux virtual machine on Microsoft Azure to avoid issues with your laptop, etc. Please remember this can be any Linux /Unix machine including your laptop. It is simply an environment you will use to deploy a cluster on Microsoft Azure.

****In the case we provide you with a virtual machine to login, you will skip steps 5, 6, 7, 8 below and start with Azure subscription or Task 2 - Setup Microsoft Azure Subscription ****

Task 1 - Deploy software on an Ubuntu Linux Server

1. On Microsoft Azure Management Portal, click **New** -> **Compute** -> **Virtual Machine** -> **Quick Create**, input all required fields including *DNS Name*, *Image* = **Ubuntu Server 12.04 LTS**, *Size* = Small, *Password* and *Region*:

NEW

COMPUTE

DATA SERVICES

APP SERVICES

NETWORK SERVICES

STORE PREVIEW

WEB SITE

VIRTUAL MACHINE

MOBILE SERVICE

CLOUD SERVICE

QUICK CREATE

FROM GALLERY

DNS NAME

.cloudapp.net

IMAGE

SIZE

Ubuntu Server 12.04 L1

Small (1 core, 1.75 GB)

USER NAME: azureuser

NEW PASSWORD

CONFIRM

REGION/AFFINITY GROUP

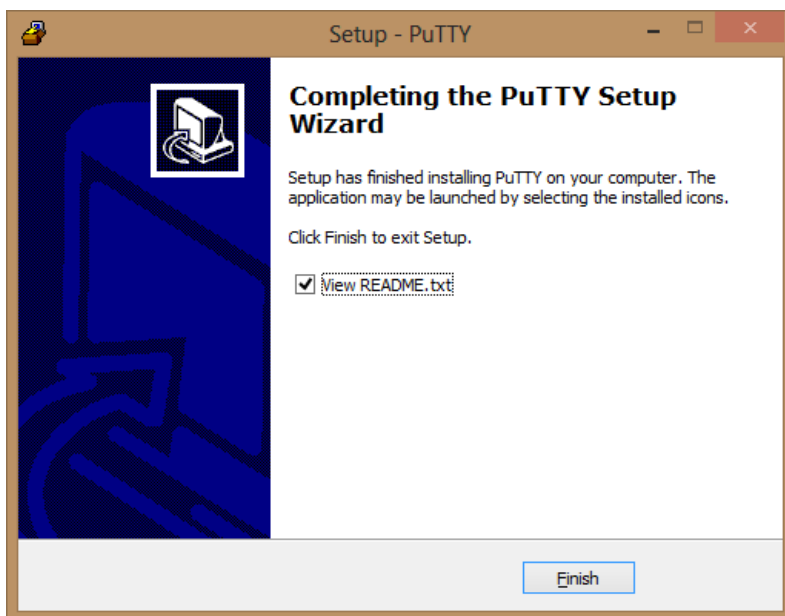
East Asia

By clicking Create, I agree this software is from Canonical and Canonical's [legal terms](#) apply to it. Microsoft doesn't provide rights for third party software.

CREATE A VIRTUAL MACHINE

Create an Ubuntu VM


- After the machine is created, we can use [putty](#) to connect to that machine from the DNS name. You can find the DNS name from the dashboard of the virtual machine. You can also install the full package of Putty if you are using windows from [Windows installer](#) for everything except PuTTYtel.



Install Putty


If you are using Linux or mac, you can directly use your own ssh tools.

- Launch Putty and connect to the remote machine with the DNS name. You can find the DNS name from the azure management portal.



ipythonmgr

DASHBOARD MONITOR ENDPOINTS CONFIGURE



web endpoint status PREVIEW

You have not configured a web endpoint for monitoring. Configure one to get started.

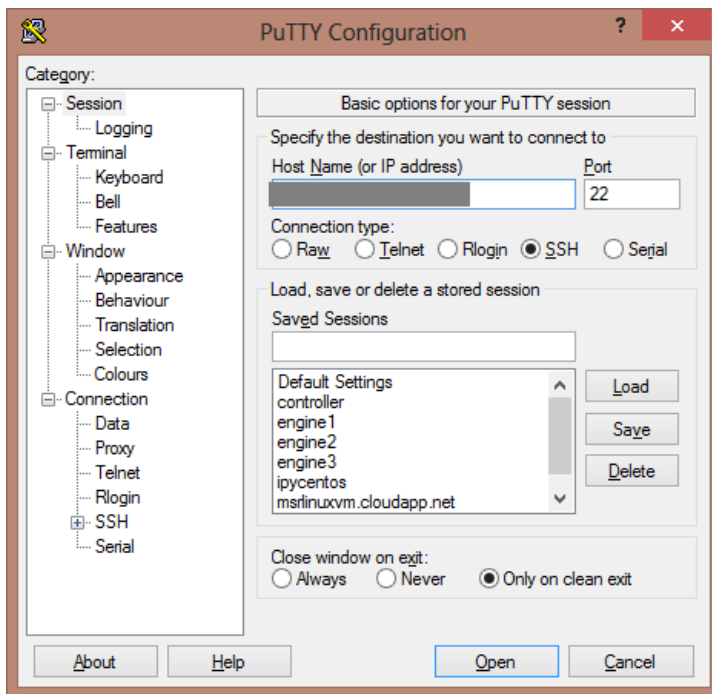
CONFIGURE WEB ENDPOINT MONITORING →

autoscale status PREVIEW

quick glance

STATUS
Running

DNS NAME
ipythonmgr.cloudapp.net



```
azureuser@ipythonmgr: ~  
login as: azureuser  
azureuser@ipythonmgr.cloudapp.net's password:  
Welcome to Ubuntu 12.04.3 LTS (GNU/Linux 3.2.0-54-virtual x86_64)  
  
* Documentation:  https://help.ubuntu.com/  
  
System information as of Tue Oct  8 03:11:17 UTC 2013  
  
System load:  0.03          Processes:            83  
Usage of /:   3.0% of 28.83GB Users logged in:      0  
Memory usage: 3%           IP address for eth0: 10.146.182.47  
Swap usage:   0%  
  
Graph this data and manage this system at https://landscape.canonical.com/  
  
Get cloud support with Ubuntu Advantage Cloud Guest:  
http://www.ubuntu.com/business/services/cloud  
  
Use Juju to deploy your cloud instances and workloads:  
https://juju.ubuntu.com/#cloud-precise  
  
0 packages can be updated.  
0 updates are security updates.  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
  
azureuser@ipythonmgr:~$
```

Connect Remote Machine

4. Then we can execute following command to install git

```
sudo apt-get install git
```

```

azureuser@ipythonmgr:~$ sudo apt-get install git
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following extra packages will be installed:
  git-man liberror-perl
Suggested packages:
  git-daemon-run git-daemon-sysvinit git-doc git-el git-arch git-cvs git-svn
  git-email git-gui gitk gitweb
The following NEW packages will be installed:
  git git-man liberror-perl
0 upgraded, 3 newly installed, 0 to remove and 0 not upgraded.
Need to get 6,741 kB of archives.
After this operation, 15.2 MB of additional disk space will be used.
Do you want to continue [Y/n]? Y
Get:1 http://azure.archive.ubuntu.com/ubuntu/ precise/main liberror-perl all 0.1
7-1 [23.8 kB]
Get:2 http://azure.archive.ubuntu.com/ubuntu/ precise/main git-man all 1:1.7.9.5
-1 [630 kB]
Get:3 http://azure.archive.ubuntu.com/ubuntu/ precise/main git amd64 1:1.7.9.5-1
 [6,087 kB]
Fetched 6,741 kB in 0s (19.9 MB/s)
Selecting previously unselected package liberror-perl.
(Reading database ... 50831 files and directories currently installed.)
Unpacking liberror-perl (from .../liberror-perl_0.17-1_all.deb) ...
Selecting previously unselected package git-man.
Unpacking git-man (from .../git-man_1%3a1.7.9.5-1_all.deb) ...
Selecting previously unselected package git.
Unpacking git (from .../git_1%3a1.7.9.5-1_amd64.deb) ...
Processing triggers for man-db ...
Setting up liberror-perl (0.17-1) ...
Setting up git-man (1:1.7.9.5-1) ...
Setting up git (1:1.7.9.5-1) ...

```

Install Git

5. Execute following command to install Python Setup Tools.

```
sudo apt-get install python-setuptools
```

```

azureuser@ipythonmgr:~$ sudo apt-get install python-setuptools
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  python-setuptools
0 upgraded, 1 newly installed, 0 to remove and 0 not upgraded.
Need to get 441 kB of archives.
After this operation, 1,068 kB of additional disk space will be used.
Get:1 http://azure.archive.ubuntu.com/ubuntu/ precise/main python-setuptools all
 0.6.24-1ubuntu1 [441 kB]
Fetched 441 kB in 0s (2,468 kB/s)
Selecting previously unselected package python-setuptools.
(Reading database ... 51480 files and directories currently installed.)
Unpacking python-setuptools (from .../python-setuptools_0.6.24-1ubuntu1_all.deb)
...
Setting up python-setuptools (0.6.24-1ubuntu1) ...

```

Install Python Setup Tools

6. Execute following commands to install Microsoft Azure SDK for Python

```

git clone https://github.com/WindowsAzure/azure-sdk-for-python.git
cd ~/azure-sdk-for-python
sudo python setup.py install

```

```

azureuser@ipythonmgr:~$ git clone https://github.com/WindowsAzure/azure-sdk-for-
python.git
Cloning into 'azure-sdk-for-python'...
remote: Counting objects: 617, done.
remote: Compressing objects: 100% (338/338), done.
remote: Total 617 (delta 314), reused 557 (delta 268)
Receiving objects: 100% (617/617), 3.20 MiB | 590 KiB/s, done.
Resolving deltas: 100% (314/314), done.

```

```

byte-compiling /usr/local/lib/python2.7/dist-packages/azure/http/batchclient.py
to batchclient.pyc
byte-compiling /usr/local/lib/python2.7/dist-packages/azure/http/httpclient.py t
o httpclient.pyc
byte-compiling /usr/local/lib/python2.7/dist-packages/azure/http/winhttp.py to w
inhttp.pyc
byte-compiling /usr/local/lib/python2.7/dist-packages/azure/http/__init__.py to
__init__.pyc
byte-compiling /usr/local/lib/python2.7/dist-packages/azure/servicebus/servicebu
sservice.py to servicebuservice.pyc
byte-compiling /usr/local/lib/python2.7/dist-packages/azure/servicebus/__init__
.py to __init__.pyc
running install_egg_info
Writing /usr/local/lib/python2.7/dist-packages/azure-0.7.0.egg-info

```

Install Microsoft Azure SDK for Python

7. Install WorkerPool and Paramiko

```

sudo easy_install workerpool
sudo easy_install paramiko

```

```

Installed /usr/local/lib/python2.7/dist-packages/workerpool-0.9.2-py2.7.egg
Processing dependencies for workerpool
Finished processing dependencies for workerpool
azureuser@ipythonmgr:~/azure-sdk-for-python/src$ sudo easy_install paramiko
Searching for paramiko
Best match: paramiko 1.7.7.1
Adding paramiko 1.7.7.1 to easy-install.pth file

Using /usr/lib/python2.7/dist-packages
Processing dependencies for paramiko
Finished processing dependencies for paramiko

```

Install WorkerPool and Paramiko

All software has been installed on your machine, next we setup to connect to Microsoft Azure Portal by Python.

Task 2 - Setup Microsoft Azure Subscription

1. To connect to the Service Management endpoint, you need your Microsoft Azure subscription ID and the path to a valid management certificate. You can obtain your subscription ID through the [management portal](#), and you can create management certificates in a number of ways. Here we use OpenSSL to create it.

You actually need to create two certificates, one for the server (a .cer file) and one for the client (a .pem file). To create the .pem file, execute this:

```

openssl req -x509 -nodes -days 365 -newkey rsa:1024 -keyout mycert.pem -out mycert.pem

```

```

azureuser@ipythonmgr:~$ openssl req -x509 -nodes -days 365 -newkey rsa:1024 -key
out mycert.pem -out mycert.pem
Generating a 1024 bit RSA private key
.....++++++
.....++++++
writing new private key to 'mycert.pem'
-----
You are about to be asked to enter information that will be incorporated
into your certificate request.
What you are about to enter is what is called a Distinguished Name or a DN.
There are quite a few fields but you can leave some blank
For some fields there will be a default value,
If you enter '.', the field will be left blank.
-----
Country Name (2 letter code) [AU]:CN
State or Province Name (full name) [Some-State]:Shanghai
Locality Name (eg, city) []:Shanghai
Organization Name (eg, company) [Internet Widgits Pty Ltd]:Training
Organizational Unit Name (eg, section) []:Azure
Common Name (e.g. server FQDN or YOUR name) []:training.com
Email Address []:training@azure.com

```

Create Client Certificate File(a .pem file)

Please input all information that required to avoid cert error problem.

- To create the .cer certificate, execute this:

```
openssl x509 -inform pem -in mycert.pem -outform der -out mycert.cer
```

You can use ls to view the file you created.

```
azureuser@ipythonmgr:~$ openssl x509 -inform pem -in mycert.pem -outform der -out mycert.cer
azureuser@ipythonmgr:~$ ls
azure-sdk-for-python  mycert.cer  mycert.pem
```

Create Server Certificate File(a .cer file)

For more information about Microsoft Azure certificates, see [Managing Certificates in Microsoft Azure](#). For a complete description of OpenSSL parameters, see the documentation at <http://www.openssl.org/docs/apps/openssl.html>.

- We need to download the cert file from the remote machine to your desktop machine. If you are using windows, you need to use pscp to download the **mycert.cer** file. Execute the following command in your windows command line:

```
cd [Your Putty Folder]
[Your Putty Folder]\pscp.exe -p [username]@[DNSName]:/home/[username]/mycert.cer mycert.cer
[Your Putty Folder]\pscp.exe -p [username]@[DNSName]:/home/[username]/mycert.pem mycert.pem
```


Just replace [Your Putty Folder],[username] and [DNSName].

```
c:\Program Files (x86)\PuTTY>pscp.exe -p azureuser@ipythonmgr.cloudapp.net:/home/azureuser/mycert.cer mycert.cer
azureuser@ipythonmgr.cloudapp.net's password:
mycert.cer          ! 0 kB !    0.7 kB/s ! ETA: 00:00:00 ! 100%
```

Download cer file

If you get an error about "Cannot create file xxx", please run your command line in administrators mode and run again.


- After you have downloaded the file *mycert.cer*, you will need to upload the .cer file to Microsoft Azure via the "Upload" button of the "Settings" tab of the management portal, and you will need to make note of where you saved the .pem file.



Upload a management certificate

Upload a certificate (.cer) file for managing your subscription.

FILE

 mycert.cer





Uploading the following management certificate for the subscription: {0}...

DETAILS



Upload .cer file to Microsoft Azure

- Please also note the subscription id for future use.

settings

MANAGEMENT CERTIFICATES ADMINISTRATORS AFFINITY GROUPS USAGE

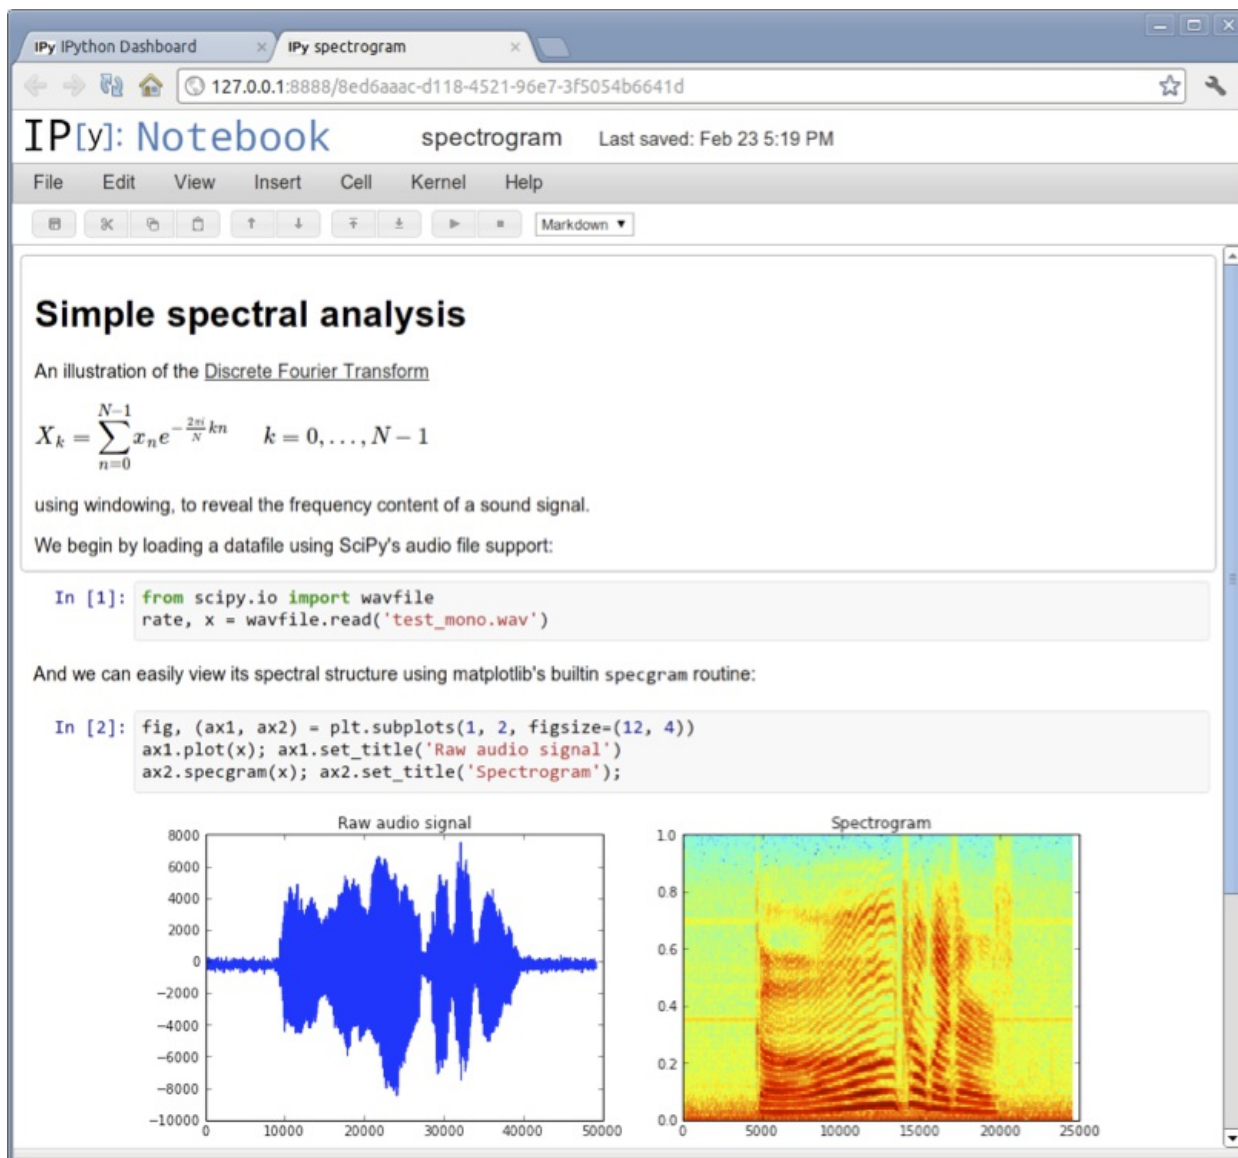
NAME	STATUS	SUBSCRIPTION	SUBSCRIPTION ID	THUMBPRINT	EXPIRES ON		
Trainer	✓ Created	Azpas300A0U8655	C11E4D9AE6A024F481982...	C1BE4D9AE6A024F481982...	10/3/2014		

Get Subscription Id

Exercise 2 - Deploy IPython Notebook on Microsoft Azure

The [IPython project](#) provides a collection of tools for scientific computing that include powerful interactive shells, high-performance and easy to use parallel libraries and a web-based environment called the IPython Notebook. The Notebook provides a working environment for interactive computing that combines code execution with the creation of a live computational document. These notebook files can contain arbitrary text, mathematical formulas, input code, results, graphics, videos and any other kind of media that a modern web browser is capable of displaying.

Whether you're absolutely new to Python and want to learn it in a fun, interactive environment or do some serious parallel/technical computing, the IPython Notebook is a great choice. As an illustration of its capabilities, the following screenshot shows the IPython Notebook being used, in combination with the SciPy and matplotlib packages, to analyze the structure of a sound recording:



IPython Notebook Spectral

1. First we need to copy the toolkit under **Source\Ex02-Deploy\Python** to local Ubuntu machine, then copy the **mycert.pem** to the same folder of the source.
2. Open the file **configSample.py** in gedit. You need to replace the subscription id with yours and the path to your private key file (**mycert.pem**).

```
#####
#           Azure Settings           #
#####
# the subscription id, replace <subscription_id> to your real id
subscription_id = "<subscription id>"
# the pem file path, replace <pem_file_name> to your real pem file name
pem_path = "<pem_file_name>"

# the number of vms
# one controller and n - 1 engines
# default 2
num_vm = 2

# service_name should be globally unique
# service_name should only contain letters, numbers and hyphen, the length should be less th
# suggest: service_name, deployment_name and role_name set to the same
service_name = "<service name>"
deployment_name = "<deployment name>"
role_name = "<role name>"

# Please replace the <storage_account_name> to real storage account name in your subscriptio

media_link_base = "http://<storage_account_name>.blob.core.windows.net/vhds/ipythonvm"
# location of the vm
location = "<your region>"
```

Edit configSample.py File

You may also need to change the following sections in that files including, **The number of nodes**, **The name of VM nodes**, **Service Name**, **Deployment Name**, **Role Name**, **Media Link Base**, **The password of the Notebook**.

For **Service Name**, **Deployment Name**, **Role Name**, ensure they are global unique. You can add your own name and number to avoid conflicts.

For the "Media Link Base", you must set in **http://.blob.core.windows.net/vhds/ipythonvm** to your correct storage account. If you don't have any storage account under your subscription, you can just create one. Replace to your real region like "East Asia" or "East US". Now we set the number of nodes to 3, which means 1 controller and 2 engines.

The screenshot shows the Azure portal interface for creating a new storage account. On the left, there's a navigation pane with categories like COMPUTE, DATA SERVICES, APP SERVICES, NETWORK SERVICES, and STORE. The main area is titled 'NEW' and has tabs for COMPUTE, SQL DATABASE, and QUICK CREATE. The 'QUICK CREATE' tab is active, showing a form to create a storage account. The 'URL' field is highlighted with a red box and contains a placeholder ending in '.core.windows.net'. The 'LOCATION/AFFINITY GROUP' dropdown is set to 'East Asia'. The 'Enable Geo-Replication' checkbox is checked. The 'CREATE STORAGE ACCOUNT' button is at the bottom right.

Create Storage Account

3. Next we need to upload all files under **Source\Ex02-DeployPython** to remote ubuntu machine.

We create a folder on that ubuntu machine in Putty.

```
cd ~
mkdir Ex02-DeployIPython
```

```
azureuser@ipythonmgr:~$ cd ~
azureuser@ipythonmgr:~$ mkdir Ex02-DeployIPython
```

Create Folder

Then we run the following commands in Windows command line locally:

```
"[You Putty Folder]\pscp.exe" -r "[Source\Ex02-DeployIPython]\*" [username]@[DNSName]:/home/[username]/Ex02-DeployIPython
```

```
d:\Ex02-DeployIPython>"c:\Program Files (x86)\PuTTY\pscp.exe" -r d:\Ex02-DeployIPython * azureuser@ipythonmgr.cloudapp.net:/home/azureuser/Ex02-DeployIPython
azureuser@ipythonmgr.cloudapp.net's password:
configSample.py      | 7 kB | 7.8 kB/s | ETA: 00:00:00 | 100%
configSample.pyc     | 6 kB | 6.3 kB/s | ETA: 00:00:00 | 100%
createMultiUM.py     | 3 kB | 3.4 kB/s | ETA: 00:00:00 | 100%
creatUM.py           | 12 kB | 12.4 kB/s | ETA: 00:00:00 | 100%
creatUM.pyc          | 11 kB | 11.1 kB/s | ETA: 00:00:00 | 100%
installMultiSoftware.py | 1 kB | 1.3 kB/s | ETA: 00:00:00 | 100%
installSoftware.py   | 4 kB | 4.3 kB/s | ETA: 00:00:00 | 100%
main.py              | 1 kB | 1.5 kB/s | ETA: 00:00:00 | 100%
mycert.cer           | 0 kB | 0.7 kB/s | ETA: 00:00:00 | 100%
mycert.pem           | 1 kB | 1.9 kB/s | ETA: 00:00:00 | 100%
ssh.py               | 7 kB | 7.7 kB/s | ETA: 00:00:00 | 100%
ssh.pyc              | 7 kB | 7.9 kB/s | ETA: 00:00:00 | 100%
util.py              | 0 kB | 0.3 kB/s | ETA: 00:00:00 | 100%
util.pyc             | 0 kB | 0.6 kB/s | ETA: 00:00:00 | 100%
configSample.py      | 7 kB | 7.8 kB/s | ETA: 00:00:00 | 100%
configSample.pyc     | 6 kB | 6.3 kB/s | ETA: 00:00:00 | 100%
createMultiUM.py     | 3 kB | 3.4 kB/s | ETA: 00:00:00 | 100%
creatUM.py           | 12 kB | 12.4 kB/s | ETA: 00:00:00 | 100%
creatUM.pyc          | 11 kB | 11.1 kB/s | ETA: 00:00:00 | 100%
installMultiSoftware.py | 1 kB | 1.3 kB/s | ETA: 00:00:00 | 100%
installSoftware.py   | 4 kB | 4.3 kB/s | ETA: 00:00:00 | 100%
main.py              | 1 kB | 1.5 kB/s | ETA: 00:00:00 | 100%
mycert.cer           | 0 kB | 0.7 kB/s | ETA: 00:00:00 | 100%
mycert.pem           | 1 kB | 1.9 kB/s | ETA: 00:00:00 | 100%
ssh.py               | 7 kB | 7.7 kB/s | ETA: 00:00:00 | 100%
ssh.pyc              | 7 kB | 7.9 kB/s | ETA: 00:00:00 | 100%
util.py              | 0 kB | 0.3 kB/s | ETA: 00:00:00 | 100%
util.pyc             | 0 kB | 0.6 kB/s | ETA: 00:00:00 | 100%
```

Copy Files

Don't forget to replace [You Putty Folder],[Source\Ex02-DeployIPython],[username],[DNSName] with your own values.

4. Execute **main.py** with following command:

```
python main.py [start|create|deploy|delete]
```

start creates VMs and deploy IPython; **create** just creates VMs; **deploy** deploys IPython on existing VMs depending on 'create'; **delete** removes all resources on Microsoft Azure.

Now we execute with **start**.

```
python main.py start
```

```
azureuser@ipythonmgr:~/Ex02-DeployIPython$ python main.py start
```

Execute Python Commands

At first, the **start** command creates a cloud service with the **Service Name** you defined in the configuration file. Then it creates 3 small instances in the cloud service. After those machines are launched, the code will connect to those machines and deploy required software and IPython Notebook automatically.

5. After about 10 minutes, the deployment is done. You will see the IPython cloud service is running in Microsoft Azure. There are 3 small instances running.

```
azureuser@ipythonmgr: ~/Ex02-DeployIPython
OK..... openssl req -new -newkey rsa:4096 -days 365 -nodes -x509 -subj /
C=CN/ST=SH/L=STAR/O=Dis/CN=pythoncluer2.cloudapp.net -keyout /home/azure/.ipytho
n/profile_nbserver/azure.pem -out /home/azure/.ipython/profile_nbserver/azure.pe
m
OK..... python -c 'from IPython.lib import passwd; print passwd("Test12"
)'
OK..... hostname -I
OK..... hostname -I
pythoncluer2.cloudapp.net 12201
OK..... ipython profile create nbclient
OK..... scp -o 'StrictHostKeyChecking no' azure@
pythoncluer2.cloudapp.net:/home/azure/.ipython/profile_nbserver/security/ipcontr
oller-engine.json /home/azure/.ipython/profile_nbclient/security/
pythoncluer2.cloudapp.net 12202
OK..... ipython profile create nbclient
OK..... scp -o 'StrictHostKeyChecking no' azure@
pythoncluer2.cloudapp.net:/home/azure/.ipython/profile_nbserver/security/ipcontr
oller-engine.json /home/azure/.ipython/profile_nbclient/security/

It's ready on: https://pythoncluer2.cloudapp.net:9090
Happy IPython!
azureuser@ipythonmgr:~/Ex02-DeployIPython$
```

pythoncluer2

→

✓ Created

✓ Running

-

East Asia

<http://pythoncluer2.cloudapp.net>


pythoncluer2

[DASHBOARD](#) [MONITOR](#) [SCALE PREVIEW](#) [INSTANCES](#) [LINKED RESOURCES](#) [CERTIFICATES](#)

NAME	STATUS	SIZE	UPDATE DOMAIN	FAULT DOMAIN	
pythoncluster20	✓ Running	Small	0	0	
pythoncluster21	✓ Running	Small	0	0	
pythoncluster22	✓ Running	Small	0	0	

IPython Notebook is running on Microsoft Azure

Just click the link on Microsoft Azure and you will see IPython is ready. If you see warnings for certification issue, just ignore it and continue.



IPython Notebook

Exercise 3 - Run Monte Carlo Simulation on IPython

Monte Carlo simulation is a computerized mathematical technique that allows people to account for risk in quantitative analysis and decision making. The technique is used by professionals in such widely disparate fields as finance, project management, energy, manufacturing, engineering, research and development, insurance, oil & gas, transportation, and the environment.

Monte Carlo simulation furnishes the decision-maker with a range of possible outcomes and the probabilities they will occur for any choice of action.. It shows the extreme possibilities—the outcomes of going for broke and for the most conservative decision—along with all possible consequences for middle-of-the-road decisions.

The technique was first used by scientists working on the atom bomb; it was named for Monte Carlo, the Monaco resort town renowned for its casinos. Since its introduction in World War II, Monte Carlo simulation has been used to model a variety of physical and conceptual systems.

In this exercise, you will run a Monte Carlo simulation code in your IPython notebook. This notebook shows how to use IPython.parallel to do Monte-Carlo options pricing in parallel. We will compute the price of a large number of options for different strike prices and volatilities, where each task will consist of computing the option price for a single strike price and volatility.

1. Login your IPython Notebook with your predefined password, it is **Test12** if you didn't change it.

IP[y]: Notebook

Logout

NotebooksClusters

To import a notebook, drag the file onto the listing below or [click here](#).

RefreshNew Notebook

/ home / azure /

Notebook list empty.

Login IPython Notebook

2. Create a new notebook. Execute the following command in a cell.

```
!wget https://raw.githubusercontent.com/wenming/BigDataSamples/master/ipythonMLSamples/Cluster%20-%20ParallelMCOptions.ipynb
```

In [2]:

```
!wget https://raw.githubusercontent.com/wenming/BigDataSamples/master/ipythonMLSamples/Cluster%20-%20ParallelMCOptions.ipynb

--2013-09-13 12:54:10-- https://raw.githubusercontent.com/wenming/BigDataSamples/master/ipythonMLSamples/Cluster%20-%20ParallelMCOptions.ipynb
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 103.245.222.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|103.245.222.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 80446 (79K) [text/plain]
Saving to: 'Cluster - ParallelMCOptions.ipynb'

100%[=====] 80,446 382K/s in 0.2s

2013-09-13 12:54:11 (382 KB/s) - 'Cluster - ParallelMCOptions.ipynb' saved [80446/80446]
```

Load Monte Carlo Simulation

3. Click the Notebooks tab. There is a new notebook called **ParallelMCOptions-cluster**, click the notebook.

NotebooksClusters

To import a notebook, drag the file onto the listing below or [click here](#).

RefreshNew Notebook

/ home / azure /

Cluster - ParallelMCOptionsShutdown

Untitled0Shutdown

Parallel Monto-Carlo options pricing with IPython

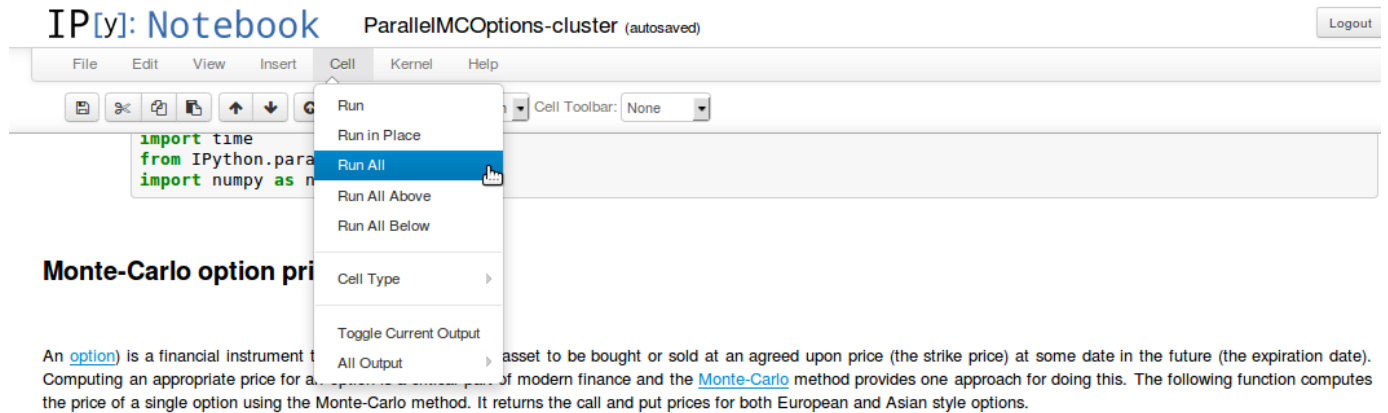
In addition to having an interactive console and Notebook web application, IPython has a powerful and high-performance parallel computing framework. This framework is fully integrated with the IPython Notebook, which opens the door for seamless parallel computing, accessible through the browser.

The fundamental unit of work in IPython's parallel computing framework (`IPython.parallel`) is the Python function. Once an IPython Cluster has been started, the API of `IPython.parallel` allows Python functions to be scheduled, along with their arguments, to be run in parallel on the Cluster. `IPython.parallel` supports a wide range of scheduling options; this tutorial will illustrate a dynamic load balancing scheduling algorithm.

This notebook shows how to use `IPython.parallel` to do Monte-Carlo options pricing in parallel. We will compute the price of a large number of options for different strike prices and volatilities, where each task will consist of computing the option price for a single strike price and volatility.

Cluster - ParallelMCOptions

- Click **Cell->Run All** to execute the sample. It will run the simulation in parallel.



Execute Monte Carlo Simulation in Parallel on Azure

The IPython Notebook and `IPython.parallel` enable you to parallelize your code on a remote cluster using nothing more than a web browser. As this example shows, once you have a Python function that performs a unit of work, it is easy to invoke that function in parallel for different arguments. The example shown here is extremely simple; the full API is rich and powerful. Details can be found in the [IPython Documentation](#).

After couple of minutes, you will see the result in the page.

```
In [10]:
%%timeit -n1 -r1

global async_results
async_results = []
for strike in strike_vals:
    for sigma in sigma_vals:
        # This line submits the tasks for parallel computation.
        ar = view.apply_async(price_option, price, strike, sigma, rate, days, paths)
        async_results.append(ar)

c.wait(async_results) # Wait until all tasks are done.
```

1 loops, best of 1: 2min 56s per loop


```
In [11]:
len(async_results)
```

Out
[11]:

400

Out
[17]:

<matplotlib.text.Text at 0x3c53410>



Monte Carlo Simulation Result

5. We can use SSH to controller and engines to check the CPU load for different machines. Run *top* command to check the most CPU consumption process.

top

During the Monte Carlo is executing, the ipengine consumes almost 100% CPU on all engine machines.

```
azure@azurecomputer: ~  
top - 17:21:01 up 20 min, 1 user, load average: 0.64, 0.58, 0.53  
Tasks: 90 total, 2 running, 88 sleeping, 0 stopped, 0 zombie  
Cpu(s): 2.8%us, 0.7%sy, 0.0%ni, 96.1%id, 0.0%wa, 0.0%hi, 0.4%si, 0.0%st  
Mem: 1720536k total, 1199652k used, 520884k free, 47260k buffers  
Swap: 0k total, 0k used, 0k free, 852040k cached  
  
  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND  
10388 azure    20   0 602m  46m 8500 R  5.0   2.8   0:03.74 python  
9902 azure    20   0 254m  29m 6284 S  2.0   1.7   0:01.83 ipcontroller  
9912 azure    20   0 257m  29m 2320 S  1.3   1.8   0:01.17 ipcontroller  
9909 azure    20   0 253m  24m 1416 S  0.3   1.4   0:00.06 ipcontroller
```

```
azure@azurecomputer: ~  
top - 17:23:58 up 22 min, 1 user, load average: 0.64, 0.50, 0.50  
Tasks: 83 total, 2 running, 81 sleeping, 0 stopped, 0 zombie  
Cpu(s): 99.0%us, 0.3%sy, 0.0%ni, 0.7%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st  
Mem: 1720536k total, 1060956k used, 659580k free, 46616k buffers  
Swap: 0k total, 0k used, 0k free, 851932k cached  
  
  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND  
9892 azure    20   0 414m  31m 7108 R 98.6   1.9   1:40.72 ipengine  
  1 root        20   0 24336 2220 1272 S  0.0   0.1   0:00.95 init  
  2 root        20   0      0      0      0 S  0.0   0.0   0:00.00 kthreadd  
  3 root        20   0      0      0      0 S  0.0   0.0   0:00.08 ksoftirqd/0  
  6 root        RT    0      0      0      0 S  0.0   0.0   0:00.00 migration/0  
  7 root        RT    0      0      0      0 S  0.0   0.0   0:00.00 watchdog/0  
  8 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 cpuset  
  9 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 khelper  
10 root        20   0      0      0      0 S  0.0   0.0   0:00.00 kdevtmpfs  
11 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 netns  
12 root        20   0      0      0      0 S  0.0   0.0   0:00.00 sync_supers  
13 root        20   0      0      0      0 S  0.0   0.0   0:00.00 bdi-default  
14 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 kintegrityd  
15 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 kblockd  
16 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 ata_sff  
17 root        20   0      0      0      0 S  0.0   0.0   0:00.00 khubd  
18 root         0 -20      0      0      0 S  0.0   0.0   0:00.00 md
```



```
azure@azurecomputer: ~  
top - 17:24:00 up 21 min, 1 user, load average: 0.60, 0.47, 0.49  
Tasks: 83 total, 2 running, 81 sleeping, 0 stopped, 0 zombie  
Cpu(s): 98.7%us, 0.0%sy, 0.0%ni, 1.3%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st  
Mem: 1720536k total, 1061072k used, 659464k free, 47040k buffers  
Swap: 0k total, 0k used, 0k free, 851764k cached  


| PID   | USER  | PR | NI  | VIRT  | RES  | SHR  | S | %CPU | %MEM | TIME+   | COMMAND     |
|-------|-------|----|-----|-------|------|------|---|------|------|---------|-------------|
| 9913  | azure | 20 | 0   | 414m  | 31m  | 7108 | R | 98.9 | 1.9  | 1:42.64 | ipengine    |
| 10114 | azure | 20 | 0   | 17336 | 1200 | 896  | R | 0.3  | 0.1  | 0:00.35 | top         |
| 1     | root  | 20 | 0   | 24408 | 2224 | 1272 | S | 0.0  | 0.1  | 0:00.94 | init        |
| 2     | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | kthreadd    |
| 3     | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.09 | ksoftirqd/0 |
| 5     | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.32 | kworker/u:0 |
| 6     | root  | RT | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | migration/0 |
| 7     | root  | RT | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | watchdog/0  |
| 8     | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | cpuset      |
| 9     | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | khelper     |
| 10    | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | kdevtmpfs   |
| 11    | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | netns       |
| 12    | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | sync_supers |
| 13    | root  | 20 | 0   | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | bdi-default |
| 14    | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | kintegrityd |
| 15    | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | kblockd     |
| 16    | root  | 0  | -20 | 0     | 0    | 0    | S | 0.0  | 0.0  | 0:00.00 | ata_sff     |


```

IPython Cluster CPU consumption

Summary

By completing this hands-on lab you learned the following:

- Provision Virtual machines with Python.
- Deploy IPython notebook on your virtual machines.
- Run Monte Carlo Simulation on IPython in parallel.

If you would like to read more about IPython cluster: http://ipython.org/ipython-doc/dev/parallel/parallel_process.html

Copyright 2014 Microsoft Corporation. All rights reserved. Except where otherwise noted, these materials are licensed under the terms of the Apache License, Version 2.0. You may use it according to the license as is most appropriate for your project on a case-by-case basis. The terms of this license can be found in <http://www.apache.org/licenses/LICENSE-2.0>.